

ML Project 2 - Report

Sarah Mikami, Samuel Waridel and Florian Kolly
CS433, EPFL 2024 ML4Science project — TNE Laboratory

Abstract—

I. INTRODUCTION

Machine Learning (ML) has become a central tool in neuroscience, with applications in neural engineering (e.g. Brain-Computer Interfaces, BCIs), diagnostics and fundamental research. For instance, Jalilifard et al. [1] demonstrated the efficacy of Support Vector Machine (SVM) classifiers in basic emotions recognition, while Bose et al. [2] used Random Forest classifiers to automate seizure detection in epileptic patients with high accuracy. Indeed, all major ML methods have been applied to neural signals classification[3].

In this work, we examine the applications of such ML techniques in three classification problems using sEEG signals: i) distinguishing between the execution and observation of a movement; ii) differentiating between grasping a small versus a large object; iii) differentiating between observing the grasp of a small versus a large object.

Section II outlines the data acquisition process and its relevance to our classification tasks. In Section III, we describe the preprocessing of the sEEG signals and the extraction of relevant features. Sections IV and V present the models applied to our classification problems and their performance. In Section VI, we explore the idea of using only Deep Learning (DL) to directly extract features from the EEG signal. Finally, we discuss our results in Section VII and conclude in Section VIII.

II. DATA ACQUISITION AND PROBLEM SPECIFICATION

The dataset comprises signals recorded from four participants, each implanted with sEEG electrodes in distinct brain regions. During multiple experimental sessions, participants were instructed to either observe or execute specific motor tasks. These tasks involved two types of movements: a palmar grasp (force-based movement) and a pinch grasp (precision-based movement). Each session included multiple trials structured as follows: in an execution trial, the trial begins with a start cue. One second later, a light cue signals the type of object (small or large) to be targeted. This light remains illuminated for 0.5 seconds before turning off. After an additional 1.5 seconds, a "go" signal prompts the participant to lift their hand from the resting position, grasp the indicated object, return it to its original location, and place their hand back in the resting position. Observation trials followed an identical structure, except that participants passively observed the experimenter performing the actions.

A light cue indicated the participant's role (execution or observation) in each trial.




Figure 1. Timing of the trials

III. DATA ANALYSIS

In this section, we describe the preprocessing and analysis of the data. The dataset is a dictionary containing four subjects. Each subject participated to a number of sessions, ranging from only one session (subjects s6 and s7) up to three (subject s11). During each sessions, the subjects completed 256 tasks, equally distributed as observation or execution tasks. The channels are not identical across subjects.

A. Channel responsiveness

We first filtered the channels of interest: for a given channel and task, we define i) the *baseline signal* as the average signal of all trials corresponding to the task in the one second between the start of the trial (1) and the cue light turning on (2), ii) the *effect signal* as the average signal of all trials corresponding to the task in the one second around the instant at which the object is grasped (6) (either by the subject, or the experimenter). Using Welch's method[4], we compute the power spectral density (PSD) for both the *baseline signal* and the *effect signal*. We then compare the PSDs of those signals in a t-test and define a *responsive channel* as a channel with a statistically significant t-test ($\alpha < 0.05$). This step is done simultaneously for channels that are responsive only in the observation trials, channels that are responsive only in the execution trials, and channels that are responsive for both types. As the channels are not identical across subjects, the results are specific to a single subject, so the preprocessing has to be done subject-wise.

B. Data preprocessing

The signals (one signal per channel) of each trial is preprocessed in three steps.

- 1) Subsampling: the signal is subsampled from 2048Hz to 500Hz to reduce the computation needs. As our

EEG data contains information up to around 150Hz, this subsampling respect the Nyquist criterion.

- 2) Z-score correction: Each EEG signal is normalized with Z-score normalization using the mean and standard deviation of the *baseline signal*
- 3) Resampling: finally, we resample our trials such that all have the same number of timepoints

C. Feature extraction

TODO

- Amplitude, shape of response → difference?
- Response pattern at the same timing?
- Similar frequency characteristics?
- Correlation? High correlation suggests congruence
- Are the results obtained similar for both movements?

IV. ACTION RECOGNITION

In this first classification task, we train and optimize different family of ML models to predict whether the participant was executing the movement, or observing it. We use only channels that are responsive to both movements. Due to the very different tasks compared here, we expect all our models to perform quite well.

A. Relevant channels

Which regions have the most responsive channels?

- Across participant: same region? same nb of responsive channels?
- Do we find the expected regions? Link to mvt execution / observation?

For channels where we have response for both ex & obs, are the responses congruent (similar pattern) or incongruent (different)?

B. Models

Logistic regression: Logistic regression (LR) is a computationally simple ML technique, and has previously used in EEG classification[5], [6]. Here, we use LR to create two different baselines on which to compare more complex models. Firstly, we selected non-responsive channels and extracted features from them. This first baseline, which we expect to work only slightly above chance, serves as a confirmation that our channel selection is indeed a valuable initial preprocessing step. A second baseline is to train another LR model on the features extract from the responsive channel. It is the performance of this model on which we will compare all other models. We also applied PCA to reduce the dimension of our data to create a third LR model. Every time PCA is applied, we keep 95% of the explained variance.

SVM: We train two Singular Vector Machines (SVM) models, with and without PCA. SVM can handle high dimensional data, even for small datasets. Although more computationally than other methods, it is widely used in EEG classification with good results [7].

Random forest: Random Forest (RF) classifiers combine the output of multiple decision trees to reach a single result. Decision trees are algorithms classify data by creating branches depending on the value of a feature: at the end of all the branches, the leaves provide a classification result. As decision trees are prone to errors and overfitting, it is recommended to combine multiple decision trees together to improve the accuracy. A RF classifier do exactly that, by adding additional constraints to make sure the trees are different.

MLP: In addition to traditional ML models, we also train Multilayer Perceptron (MLP). We chose to include MLP as it is widely used in EEG classification¹. The main issue with MLP is that they are universal approximators and are thus prone to overfitting, especially for datasets as small as ours.

C. Results

D. Discussion

Which freq. bands are important? Is it the same across subjects?

V. MOVEMENT RECOGNITION

For the second and third type of classification tasks, we analyze whether ML tools are able to classify the type of movement done (pinch grasp or force grasp) during execution and observation. We expect both tasks to be hard to learn. Observation, in particular, is a complex task, as the participant is simply seeing two different movements that have the same goal of lifting an object. In this part, we reduce the restrictions to select the *relevant channels*: it suffices for a channel to be relevant in execution or in observation respectfully to be included.

A. Responsive channels

Which regions have the most responsive channels, are they the same for execution & observation?

- Across participant: same region? same nb of responsive channels?
- Are the channels responsive for obs also responsive for ex? Vice-versa?
- Do we find the expected regions? Link to mvt execution / observation?

For channels where we have response for both ex & obs, are the responses congruent (similar pattern) or incongruent (different)?

B. Models

We use the same families of model as in the previous part, and this for both classifying in execution tasks and in observation tasks.

¹3'610 results for 'EEG MLP classification' on Google Scholar at the time of writing

C. Results

As expected, our results are significantly lower in those tasks than in the first part. We also observe a drop in performance between execution and observation.

D. Discussion

Which freq. bands are important? Is it the same across subjects?

VI. PURE DEEP LEARNING APPROACH

Instead of extracting features on the relevant channels before training deep networks, one can ask whether the networks should learn to do that themselves. It's this question that we explore in this last section. We train two types of Convolution Neural Networks (CNN) directly on the trial signals. These networks contain several convolutional layers followed by linear layers. The hope is that the convolutional layers manage to extract relevant features from the signal, giving it as input to the subsequent MLP.

A. 1D CNN

One-dimensional CNNs are composed of kernels that are convolved with the signal directly. As they operate on data in one dimension, the input will be a tensor of size (N, C, L) , where N is the batch size, C the number of (relevant) channels and L the fixed length of each trial. Figure 2 shows how a kernel works in such network. This is close to how we computed how features: supposing the same kernel size as the length of the moving averages we use to create features, we would get the same features if the network learned moving averages as well. Given this and the known power of CNNs, we expect this model to have the best performance. The caveat is overfitting, as we have small datasets.




Figure 2. Example of a 1D kernel (taken from [8])

B. 2D CNN

Two-dimensional CNNs are composed of 2D kernels that will operate on multiple channels at the same time. This means that we give as input a tensor of dimensions (N, C, H, W) , with H the number of channels and W the length of each signal. The use of H and W as dimensions is not an arbitrary choice. Indeed, we can think of this superposition of channels as a picture, where each signal is a position on the y -axis and every timepoint a position on the x -axis. Here, $C = 1$ as this virtual picture is simply represented by nuances of grey.

C. Results

Both type of CNN works well.

VII. DISCUSSION

VIII. CONCLUSION

IX. ETHICAL RISKS

X. APPENDIX




Figure 3. Experimental setup

ACKNOWLEDGEMENTS

We thank Leonardo Pollina for taking us with him on this project, and for his quick and helpful answers to all our questions.

REFERENCES

- [1] A. Jalilifard, E. B. Pizzolato, and M. K. Islam, "Emotion classification using single-channel scalp-eeeg recording," in *2016 38th Annual International Conference of the IEEE Engineering in Medicine and Biology Society (EMBC)*, 2016, pp. 845–849.
- [2] S. Bose, V. Rama, N. Warangal, and C. Rama Rao, "Eeg signal analysis for seizure detection using discrete wavelet transform and random forest," in *2017 International Conference on Computer and Applications (ICCA)*, 2017, pp. 369–378.
- [3] M.-P. Hosseini, A. Hosseini, and K. Ahi, "A review on machine learning for eeg signal processing in bioengineering," *IEEE Reviews in Biomedical Engineering*, vol. 14, pp. 204–218, 2021.
- [4] O. M. Solomon, Jr, "Psd computations using welch's method. [power spectral density (psd)]." 12 1991. [Online]. Available: <https://www.osti.gov/biblio/5688766>
- [5] A. Subasi and E. Erçelebi, "Classification of eeg signals using neural network and logistic regression," *Computer Methods and Programs in Biomedicine*, vol. 78, no. 2, pp. 87–99, 2005. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0169260705000246>
- [6] R. Tomioka, K. Aihara, and K.-R. Müller, "Logistic regression for single trial eeg classification," in *Advances in Neural Information Processing Systems*, B. Schölkopf, J. Platt, and T. Hoffman, Eds., vol. 19. MIT Press, 2006. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2006/file/35937e34256cf4e5b2f7da08871d2a0b-Paper.pdf

- [7] M. N. A. H. Sha'abani, N. Fuad, N. Jamal, and M. F. Ismail, "knn and svm classification for eeg: A review," in *InECCE2019*, A. N. Kasruddin Nasir, M. A. Ahmad, M. S. Najib, Y. Abdul Wahab, N. A. Othman, N. M. Abd Ghani, A. Irawan, S. Khatun, R. M. T. Raja Ismail, M. M. Saari, M. R. Daud, and A. A. Mohd Faudzi, Eds. Singapore: Springer Singapore, 2020, pp. 555–565.
- [8] A. Baldominos, Y. Sáez, and P. Isasi, "Evolutionary design of convolutional neural networks for human activity recognition in sensor-rich environments," *Sensors*, vol. 18, 04 2018.